

Lecture 1: Basics

Capture Devices

Passive:

→ RGB

→ Stereo and Multi-View

Stereo Matching: Sparse vs. Dense Matching

Wider dist. between cameras: Feature extraction vs. Nat.

↑ distortion \rightarrow occlusion \leftarrow

Smaller dist. between cameras: Search window: smaller = more noise
larger = less details

↑ error

+ work indoors + outdoors

- Need Computer

- Features needed

+ No features needed

- Fails outdoors

- Sensitive to battery

Active:

→ ToF (Time of flight)

→ Structured light

→ LIDAR

- slow

+ very precise

↓ ↓ ↓

↓ ↓ ↓

Same as ToF
but needs precise
calibration between
projector and sensor

- Correspondence Failures:
- 1) Non-Lambertian surfaces \rightarrow specular reflection
(\rightarrow not diffuse)
 - 2) Occlusions / Repetition
 - 3) Textureless Surfaces

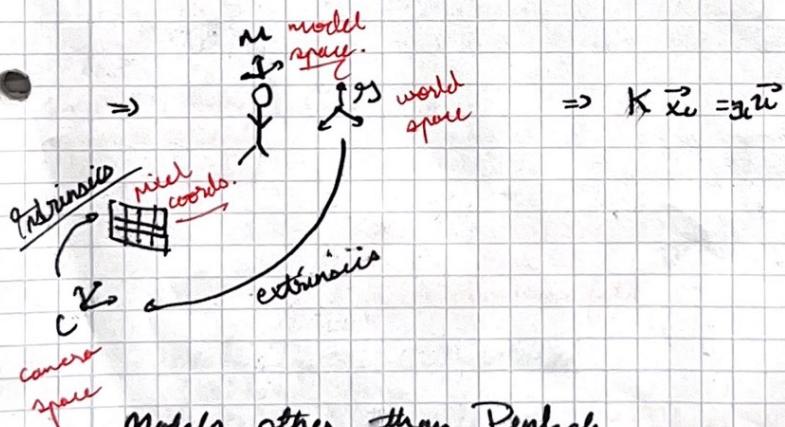
Camera Concepts

- Intrinsic: - Focal Length
- Principal Point.
- (Skew)
- (Distortion)

$$K = \begin{pmatrix} f_x & r & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{pmatrix}$$

r : shear coefficient
 c_x, c_y : principal
points or
optical center.

- Extrinsic: R, t .



$$\Rightarrow K \vec{x}_c = \vec{z} \vec{u}$$

f_x, f_y : focal
length.

Models other than Pinhole

- Radial + Tangential Distortions: Use approximate values using camera calibration.

Fish eye lens.

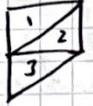
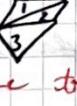
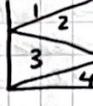
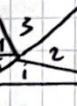
\hookrightarrow image plane not parallel to lens

Lecture 2: Surface Representations

obviously efficient
→ Polygonal Meshes:

- Error in shape representation: $O(2^k)$ i.e. doubling the # of edges gives $\times \left(\frac{1}{2}\right)^2$ reduction in error.
- Allows for recursive subdivisions
↳ i.e. higher resolution

Topology (triangulation) vs. Geometry

Topologically equivalent	=	
Geometrically Different	=	
Geometrically equivalent	=	
Topologically Different	=	

↳ same triangulation

↳ different triangulation but same rectangle.

Manifolds

* Triangle is called manifold if:

~~non manifold~~
~~self intersecting~~
~~empty interior~~
~~multiple vertices~~
~~multiple edges~~
~~multiple faces~~
~~non manifold~~

(i) Intersection of two triangles are (+ line):
 ↳ empty i.e. no area.
 ↳ common vertex
 ↳ common edge.

here intersection spans the area of small triangles = not empty.
 ↳ Non-manifold

(ii) Edges have:

- * One adjacent triangle → border edge
- * Two adjacent triangles → inner edge

edge
 ab has three
 adjacent triangle
 ↳ Non-manifold

(iii) Vertices:

- * Build a single open fan → border vertex
- * Build a single closed fan → inner vertex

vertex connected to
 two open fans
 ↳ non-manifold

vertex is connected to
 two closed fans = Non-manifold

Mesh Data Structures

Shared Vertex

+ OFF Obj STL etc. ↳ PLY allows for binary storage

Vertex data structure

+ List v → List f.

v ...
 vn-i ...
 f ...

* (ii) Half-Edge Data Structure

no only accepts manifold

representations (since stores adjacent edges)

↳ store vertices + faces + edges (i.e. opposite edge
 adjacent edge etc.)

Useful for geometric computations a)

Surfaces

(i) Explicit surfaces: $f: \mathbb{R}^2 \rightarrow \mathbb{R}$ i.e. given x, y and f give height
 (ii) Parametric surfaces:

Constructed from Bezier Curves \rightarrow Bezier patches

CAD, etc
Used in 3D design
+ construction, etc.

Bezier curve

Bezier patch:
 $f(u, v): \mathbb{R}^2 \rightarrow \mathbb{R}^3$

(iii) Constructive Solid Geometry:

Boundaries created from boolean operations:



(iv) Implicit Surfaces:

\rightarrow SDF
~~Heat~~
~~Distance~~
 \rightarrow Density

Defining Implicit Surfaces:

\rightarrow local method

or

(i) Hoppe method: $f(\vec{x}) = (\vec{x} - \vec{p}) \cdot \vec{n}_p \rightarrow$ normal of point \vec{p}
for any \vec{x} \vec{p} \rightarrow nearest \vec{p} (from point cloud e.g.)

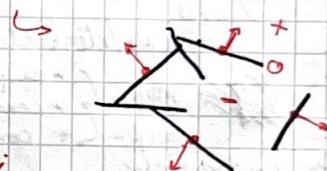
\rightarrow bad

with outliers

\rightarrow Creates Voronoi diagram which defines closest points

\rightarrow Piecewise linear

\rightarrow Discontinuous along edges of voronoi



\rightarrow global method

(ii) RBF: Radial Basis Functions:

\rightarrow Better at

handling outliers

We assume smoothness $\rightarrow f(\vec{x}) = \sum_{i=1}^n \alpha_i \varphi_i(\vec{x}) + b \cdot \vec{x} + d$

To solve: $f(\vec{p}_i) = 0$

$$\begin{aligned} f(\vec{p}_i + \varepsilon \vec{n}_p) &= \varepsilon \\ f(\vec{p}_i - \varepsilon \vec{n}_p) &= -\varepsilon \end{aligned} \quad \left. \begin{array}{l} \text{our constraints} \\ \text{some } C^2 \text{ function} \end{array} \right\} \begin{array}{l} \alpha_1 | \alpha_2 | \dots | \alpha_n \\ b \\ d \\ \hline n+3 \end{array} = n+4$$

Linear \Rightarrow SVD to solve for α

(iii) Poisson Reconstruction

$$\nabla \cdot \nabla f(\vec{x}) = \nabla \cdot \vec{N} \quad \text{normal vector field}$$

\rightarrow basically divergence of all our normals

divergence of our $\nabla f(\vec{x})$

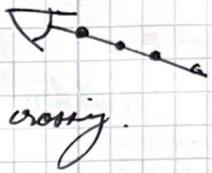
gradient field

for poison
we don't use SDF


Rendering Implicit Surfaces

(i) Ray Casting

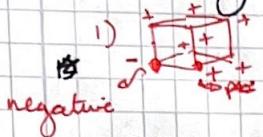
- 1) Ray Marching:
 * Apply fixed step length.
 * linearly interpolate if zero crossing.



- 2) Sphere Tracing:
 * Dynamic Steps
 * Stop if distance below threshold.
 (Go-to method usually).



- (ii) Marching Cubes:
 1) Determine zero crossings for all grid cells.



- 2) If zero-crossing use lookup table to triangulate

- 3) Linearly interpolate triangulation

negative
 as ID in Lookup table.

3) ~~Find~~ Find Triangulation

4) Linearly interpolate triangulation

Lecture 3: 3D reconstruction Methods

Feature Detection

(i) FAST Detectors: for pixel $p(x,y)$ if ≥ 12 pixels are brighter

⇒ feature points

partial derivatives w.r.t. $x, y: I(u,v) - I(u+1,v)$
etc.

in circular kernel for example

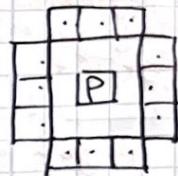
(ii) Harris Corner Detectors: if we have I_x, I_y :

we can do SSD
for e.g. Harris
but not robust.

$$H = \begin{pmatrix} \sum_w I_x^2 & \sum_w I_x I_y \\ \sum_w I_x I_y & \sum_w I_y^2 \end{pmatrix}$$

is Hessian of some image patch W . basically

we find λ_1, λ_2 of H :
 $\lambda_1 \sim \lambda_2 + \text{small} \Rightarrow \text{flat region}$
 $\lambda_1 > \lambda_2 \Leftrightarrow \text{OR } \lambda_2 > \lambda_1 \Rightarrow \text{Edge regions}$
 $\lambda_1 \sim \lambda_2 + \text{large} \Rightarrow \text{corner regions.}$



Trick to find λ_1, λ_2 w/o SVD: $R = \det(M) - \alpha (\text{trace}(M))^2$
 $= \lambda_1 \lambda_2 - \alpha (\lambda_1 + \lambda_2)^2$ properties of det + trace

if $\lambda_1 \sim \lambda_2 \sim 0 \Rightarrow R \sim 0 \Rightarrow$ flat region
if $\lambda_1 > \lambda_2 \Rightarrow \lambda_1 \lambda_2 \sim 0 \Rightarrow R = -\alpha \lambda_1^2 \Rightarrow$ edge
if $\lambda_1 \sim \lambda_2 \Rightarrow \lambda_1 \lambda_2 - \alpha \lambda_1^2 \Rightarrow$ corner.
↳ usually $0.04-0.00$

if $|R|$ small \Rightarrow flat region
 $R < 0 \Rightarrow$ edge region
 $R > 0 \Rightarrow$ corner region

{ cool trick actually :)

(iii) SIFT Detectors: finds features that are scale invariant + filter (usually Haar LS)
(And rotation invariant) i.e. applies multiple scalings + blur and finds features not affected + finds dominant rotation direction to compare features.

Correspondence Matching

(i) Simple SSD per feature: shitty + not robust

(ii) SIFT Descriptor: after SIFT filters optimal features: takes 16-sub blocks neighbourhood around feature point, rotates neighbourhood in preferential direction, calculates 8-bin orientation histogram per sub-block $\Rightarrow 8 \cdot 16 = 128$ dimensional feature vector per feature / gacss.

(iii) ConvNet: Basically siamese (shared weights) network

inputting some image/volume patch from two views and outputs 1/0 depending on match or not.

(iv) SURF, ORB etc.

Main 3D recon methods:

(i) Visual Hull Carving: If we have:

- n cameras
- we know each camera's intrinsics
- we know each camera's extrinsics
- we segment subject in each camera

Carve over voxel grid based on each camera segmentation

(ii) Structure from Motion (SfM): Given: → Set of images
→ Features + correspondences
for each image

Bundle Adjustment

Reprojection energy:

$$\text{optimizes over } \text{Err-proj}(T, X) = \sum_{i=1}^m \sum_{j=1}^n \|x_{ij} - \pi_i(T_i X_j)\|_2^2$$

Images $i \rightarrow$
feature points j in image i \rightarrow 3D world coordinate of feature point j in 3D space X_j transformed to camera frame i : $T_i X_j$, then projected to camera i .

Bundle Adjustment \leftarrow { Goal: Find 3D locations X_j
+ Find Camera poses T_i .

- * Unknowns: → 6 · (m-1) ≈ Extrinsics of all cameras except first (since identity)
→ 3 · n ≈ Feature points in 3D space.
→ (optional): Camera Intrinsic per camera.

- * Constraints: → 2 · n · jn ≈ all feature points in all images in 2D space.

$$\Rightarrow 2 \cdot n \cdot jn \geq 6(m-1) + 3n + k.$$

- (iii) Multiple View Stereo (MVS): Given: → Set of images

- * We can do frame pair matching + vote

→ Poses of images T_i
→ 3D locations of features X_j

Goal: → Dense Point Cloud.

- * Global optimization is NP-Hard

- * Plug point clouds to surface fitting e.g. nonlinear surface fitting

↳ Nonlinear Pipeline: 2D images → SfM → MVS → Reconstruction → 3D mesh.

- (iv) SLAM: ≈ smaller compute than SfM + MVS.

Frame to \rightarrow

Frame OR

Frame to Model

- * Basically real-time SfM + MVS

- * Different than Bundle Adjustment since no simultaneous offline global optimization of all poses + 3D locations.

- * Types: { 1) ORB SLAM } can also either use feature points for alignment OR directly input
RGB 2) LSD SLAM.

RGB-D { 3) KinectFusion

SLAM usually has Loop-Closure Problem

RGB-D Bundling $E_{\text{bundle}}(T) = \sum_{i,j} \sum_u \|T_i p_{iu} - T_j p_{ju}\|_2^2$ same bat for j 1) Take back-projected corr. in i, 2) Transform to 3D w/ T_i , 3) to j camera frame, 4) Project to 2D

Edepth(T) = $\sum_{i,j} \sum_u \|p_{iu} - T_i T_j^{-1} T_j p_{ju}\|_2^2$ 5) Depth Value a_j

Ecolor(T) = $\sum_{i,j} \sum_u \|\nabla I(T_i(p_{iu})) - \nabla I(T_j(T_j^{-1} T_i p_{iu}))\|_2^2$ 6) back-project p_{iu} , 7) Transform to 2D w/ T_j , 8) to original j camera frame T_j

Color gradient of original point in frame i \rightarrow Color gradient of projected point a_j in frame j

Lecture 4: Optimization Methods:

Linear Least Squares

→ Means the model is linear in parameters

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 \dots$$

here X_2 can be $x^2, \log x, \frac{1}{x}$ doesn't matter
still linear since linearly multiplying by β_2 .

* Simplified to: $Ax = b$

overdetermined usually

* Usually solved as normal equation: $A^T A x = A^T b$

Solvers: as of linear least squares

→ Iterative Solvers: conjugate

1) Preconditioned Gradient Descent. most common iterative solver.

2) Gauss-Seidel Iteration

3) Jacobi Iteration

→ Direct Solvers:

related to 1) QR, LU-Decomposition

each other

I think 3) Cholesky

Non-Linear Least Squares

* No longer linear in params: e.g. $Y = \beta_0 X^{\beta_1}$ β_1 not linearly multiplied

Solvers: we're trying to solve $\underset{x}{\text{argmin}} f(x) = \|F(x)\|_2^2$

→ First Order Solvers: Gradient Descent etc.

$$\Rightarrow x_{k+1} = x_k - t \cdot \nabla f(x_k) \text{ momentum / line search standard BS.}$$

→ Second Order Solvers:

(i) Newton's Method: $x_{k+1} = x_k - H_f(x_k)^{-1} \nabla f(x_k)$ comes from Taylor

$$\text{taylor: } f(x+t) \approx f(x) + \nabla f \cdot t + \frac{1}{2} H_f \cdot t^2$$

(ii) Gauss-Newton: Approximation of Newton:

assume
 $x_k \in \mathbb{R}^n$

$$H_f \approx 2 \tilde{J}_F^T \tilde{J}_F \Rightarrow x_{k+1} = x_k - [2 \tilde{J}_F^T \tilde{J}_F]^{-1} \nabla f(x_k)$$

\tilde{J}_F $n \times n$

x_k $n \times 1$

residuals

where $F \in \mathbb{R}^m$

linear system (triv. PP)

$$\underbrace{2 \tilde{J}_F^T \tilde{J}_F}_{A} [\underbrace{x_k - x_{k+1}}_{x}] = \underbrace{\nabla f(x_k)}_{b}$$

(CPCG)

⇒ solve with iterative linear solvers like preconditioned gradient descent

conjugate

$$(iii) \text{Lauenberg: } x_{k+1} = x_k - [2 \tilde{J}_F^T \tilde{J}_F + \lambda I]^{-1} \nabla f(x_k)$$

$\lambda \downarrow \rightarrow \text{SVD}$

$\lambda \uparrow \rightarrow \text{Gradient Descent}$

↳ Tikhonov regularization

i.e. damping factor for each

term in J s.t. $\lambda \downarrow \text{if } f(x_k) < f(x_{k+1})$

(iv) Levenberg - Marguardt (LM): *faster convergence this way than Levenberg*

$$x_{k+1} = x_k - [2\mathbf{J}_F^\top \mathbf{J}_F + \lambda \cdot \text{diag}(\mathbf{J}_F^\top \mathbf{J}_F)]^{-1} \nabla f(x_k)$$

more sophisticated Levenberg since each λ scaled by curvature of (diagonal is independent parts of curve)

(v) BFGS / LBFGS:

$$x_{k+1} = x_k - B_k \nabla f(x_k) \quad \text{Basically approximation of Hessian.}$$

$$\text{where } B_{k+1} = B_k + \alpha u u^\top + \beta v v^\top$$

In Practice we use Limited Memory BFGS (LBFGS)

Outlier Handling:

? How do we convexify our problems?

course to fine over residuals

(i) ~~RANSAC~~: basically trial & error

(ii) Lifting Schemes:

$$\text{given: } f(x) = \sum r_i(x)^2 \Rightarrow \text{frobust}(x) = \sum w_i^2 r_i(x)^2 + \text{deg} \sum (1-w_i^2)^2$$

(iii) Robust Norms: usually hard to optimize | that penalizes $w=0$
our lifting term
 (linkers are $w=1$)

⇒ Iteratively Reweighted Least Squares

$$\text{frobust}(x) = \sum w_i r_i(x)^2 \text{ and } w_i = |r_i(x)|^{p-2} \text{ for some } p\text{-norm.}$$

$$\text{assuming we wanted } \underset{x}{\operatorname{argmin}} f(x) = \underset{x}{\operatorname{argmin}} \|F(x)\|_p^p$$

Derivative Types:

(i) Numeric derivatives: $\frac{df(x)}{dx} \approx \frac{f(x+h) - f(x)}{h} \quad h \approx 1\text{-c. take very small } h.$

→ Easy to implement ⇒ good for debugging

→ Slow and numerically unstable

(ii) Automatic Differentiation: Use dual numbers: $e \neq 0$

$$\text{Example: } f(x) = x^2 \quad \left. \frac{df(x)}{dx} \right|_{x=10} ? \Rightarrow f(10+e) = 100 + 20e + e^2 \quad \begin{aligned} e^2 &= 0 \\ e' &= 0 \end{aligned}$$

$$\Rightarrow \left. \frac{df}{dx} \right|_{x=0} = \frac{20}{e}$$

(iii) Symbolic Differentiation:

Basically no numerical solutions (Maple :))

Other Solvers: 1) Inequality Constraints: Lagrange multipliers etc.

2) Gradient - Free: Monte Carlo etc.

Lecture 5: Rigid Surface Tracking + Recon \rightsquigarrow (i.e. KinectFusion revision)

Kinect Fusion Reminder

weight updates



(i) Pose Estimation via ICP:

- Basic Steps of ICP:
- 1) Selecting source points \rightsquigarrow either:
 - Use all points
 - Uniform subsampling
 - Random sampling
 - Stable sampling
 - 2) Find Correspondences
 - 3) Weigh Correspondences
 - 4) Reject Outliers
 - 5) Assign Errors metric
 - 6) Minimize errors for pose estimate

If coarse to fine registration iterate for concavity problem.

• Finding Correspondences in ICP:

only works if we have depth values

a) Nearest neighbour \rightsquigarrow FLANN etc.

b) Projective correspondences \rightsquigarrow did in da project i.e. $\pi(\pi(T\pi(p)))$

We can improve matching via only checking compatible points:

$p_s \rightarrow \pi(T\pi(\pi(T\pi(p))))$
my correspondence in other frame.

- a) Color Compatibility
- b) Normal Compatibility
- c) Pruning
- d) Reject angle and distance outliers

Errors in ICP:

has closed form linear solution \Rightarrow

a) Point to Point: $\sum_i \|KT_{pi} - q_i\|_2^2 = E(T)$

idea here

is energy lowest if $p_i \notin q_i$ on same surface w/ normal n_i

b) Point to Plane: $\sum_i \|(T_{pi} - q_i) \cdot n_i\|_2^2 = E(T)$

usually solved w/ LM if not linearized

(\hookrightarrow can be linearized in $\theta \otimes \theta$

assuming small rotations $\cos \approx 1$

\rightarrow Cholesky / SVD

* Symmetry ICP: $C_p - q = (n_p + n_q)$

(ii) SDF / TSDF: (Implicit Representations):

- Basically representation of distance to zero-crossing band.

- TSDF: just applies truncation \wedge i.e. if > 1 from zero-crossing band

$/ < -1 \Rightarrow 1 / -1$ everywhere

Volumetric Fusion: \rightsquigarrow Surfel representation: point + normal + radius
 also exists current tsdf based on our current depth value

weights or running average concept: $D_{i+1}(x) = W_i(x)D_i(x) + W_{i+1}(x)d_{i+1}(x)$

can either be 1

everywhere OR closer

to camera weight, average weight (uncertainty weighted)

tsdf value at voxel i

currently (i.e. distance to iso surface) $W_{i+1} = W_i(x) + w_{i+1}(x)$

(iii) Surface Rendering

Important to note on SDFs:

a) Raycasting: Throw multiple rays from camera.

→ Surface found from sign change + interpolate to find intersection.

=0 → surface
<0 → unknown
→ Normal found from VTSDR.

Could have another surface ~~behind~~ behind

>0 → free space.
b) Marching Cubes: Basically use Marching Cubes table based on zero-crossing.

Follow Ups to KinectFusion:

1) Extended Kinect Fusion:
+ easy to implement
+ unlimited spatial volume.
- cannot re-integrate previous regions.

2) Hierarchical Fusion:
(↳ basically uses sparse voxel octrees) → 
+ fast lookups
+ efficient storage (high local res)
- costly updates (insertions/removals)

3) Voxel Hashing:
+ extremely fast updates (insertions/removals) O(1)
+ efficient storage
- ~~slow~~ lookups slower ⇒ encoding not as fast
+ cheap streaming

Loop Closure Follow Ups:

1) Elastic Fusion:
→ Surfel based
→ Detects loop closure by localization
→ Non-rigidly warp to close loop.

2) Bundle Fusion:
→ Added sparse energy term for pose estimation
→ Includes Surface de-integration if pose had wrong estimate

Lecture 6: Deformations of Non-rigid Surface Tracing

Mesh Deformations → standard.

$$E = \lambda_{fit} E_{fit}(V) + \lambda_{reg} E_{reg}(V, X)$$

we pick non-rigid regularizers which affects behaviour of our mesh.

$$E_{fit}(V) = \sum_{i \in C} \|c_i - v_i\|_2^2$$

minimize distance between vertex and target

1-ring form:

Regularizer Types → Our goal for regularizers: → smooth global edges → Global changes preserve local detail → Should be zero in undeformed state.



i) Laplacian Surface Editing:

$$E_{reg}(V, X) = \sum_{j \in N} \|v_i - v_j - (v_i' - v_j')\|_2^2$$

our far neighborhood (1-ring) → i.e. deformation of v_i' should be very similar to 1-ring neighbours v_i''

ii) As Rigid as Possible (ARAP):

$$E_{reg}(R, V) = \sum_i \sum_{j \in N} \|v_i - v_j - R_i(v_i' - v_j')\|_2^2$$

done per face over all vertices → means that after some R edge

* we could add weight to give higher priority to closer neighbours: $w_{ij} = \frac{1}{2} (\cot \alpha_{ij} + \cot \beta_{ij})$

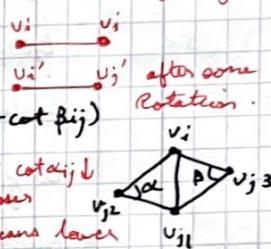
3 positions per vertex → 3 Euler angles per vertex

* Solving: Unknowns: $3n + 3n$ per position

Constraints: $3n$ valences + $3c$

average a) from E_{fit}
far neighbours

so $\cot \alpha_{ij} \uparrow \cot \beta_{ij} \downarrow$
if v_{ij} is closer then v_{ij}' means larger weight for v_{ij}'



→ We can do Levenberg Marquardt (non-linear)

OR

→ Flip-flop optimization: 1) Assume v_i, v_j' constant and solve for R

2) Assume R constant and solve for v_i, v_j'

→ Linear problem and can be solved with SVD / Procrustes.

iii) Embedded Deformation: relaxed version of ARAP

$$E_{reg}(M, V) = \sum_i \sum_{j \in N} \|v_i - v_j - M_i(v_i' - v_j')\|_2^2$$

$$E_{rot} = \sum_i \text{Rot}(M_i), \quad \text{Rot}(M) = (c_1 \cdot c_2)^2 + (c_1 \cdot c_3)^2 + \dots + (c_1 \cdot c_{i-1})^2 + \dots$$

* Solving: Unknowns: $12n$

(i.e. we want M to resemble R but allow for more DoF in our deformation (shear etc.))

Main goal:

dimensionality reduction.

⇒ quartic problem to optimize

Deformation Proxies instead of deforming mesh = computationally expensive deform proxy

we can apply: i) Cage: linear combination of vertices w.r.t. cage.

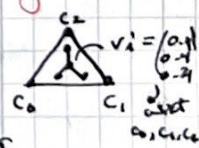
ARAP/EP

(we use barycentric coordinates for that.)

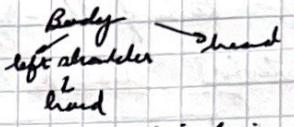
/ Igalias here.

BRUNNEN

→ we deform cages and by proxy our vertices using barycentric coordinates.



iii) 3D Grids: ~~discrete~~

e.g. Blend or iii) Skeletons: hierarchical: 
skeleton

Called linear blend skinning

{
Each vertex where is a linear combination of bone matrices: $v_i = \sum w_{i,j} M_j v_i'$
weight for each bone \rightarrow each bone has a matrix
each bone

iv) Deformation Graphs: assigns a region/neighborhood a control node.
 \rightarrow Similar to cage but instead we have deformation gone.

Non-Rigid Tracking \rightarrow find correspondences than ICP.

needs predefined mesh.

* Non-Rigid ICP: \rightarrow more degrees of freedom. e.g. ARAP/ED.

$$E = \lambda \text{point Epoint} + \lambda \text{plane Eplane} + \lambda \text{color Eccolor} + \lambda \text{try Etry}$$

from regular point to plane \rightarrow lower energy if colors match

$$\text{Epoint: } \Sigma_i \|v_i - T v_i'\|_2^2$$

$$\text{Eplane: } \Sigma_i [(v_i - T v_i') \cdot n v_i]^2 \text{ forces } v_i' \text{ to stay on the plane}$$

$$\text{Eccolor: } \Sigma_i [I(\pi(v_i)) - I'(\pi(T v_i))]^2 \text{ if } v_i \rightarrow E=0$$

$$\text{Etry: ARAP.} \quad \text{lower energy if colors match.}$$

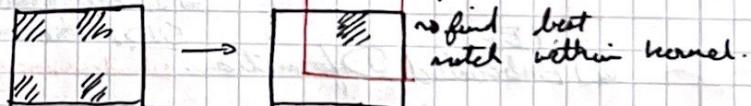
* Important hyperparameter-thresholds for non-rigid ICP:

- 1) distance
- 2) angle
- 3) Occlusion
- 4) Viewing angle

Improves with more cameras. α

88 * Correspondence Association: \rightarrow for speed.

e.g.:
1) Projective correspondences every k-th element
2) Kernel around k-th sample for minimum reduction

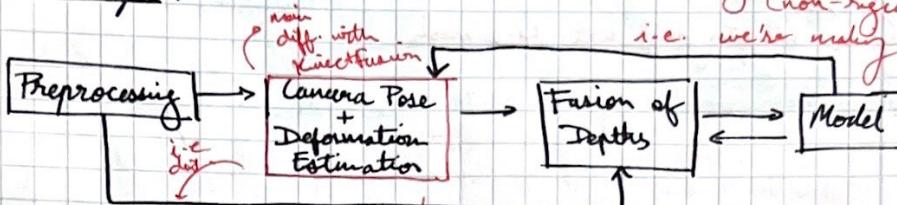


Lecture 7: Non-Rigid Tracking and Reconstruction

General Concepts:

↳ here we're tracking AND fusion
(non-rigid)

i.e. we're using mesh.



- * For Camera Pose + Deformation Estimation i.e. non-rigid ICP + ARAP etc.

- iterate
 - 1) Deform SDF
 - 2) Find 150 surface usually with marching cubes since raycasting hard for
 - 3) Find Correspondences non-rigid
 - 4) Optimize deformations non-rigid ICP with some non-rigid regularizers
 - 5) Integrate RGB-D data to model

at iteration

ICP passes

& Main Methods discussed:

In reconstruction: (i) Dynanic Fusion: cause is

we are marching cubes here.

Than this.

space noisy
on OR
skin defining?

(ii) Volume Deform:
We can increase resolution based on grid
This means higher stability compared to dynamic fusion

More stable + Better texture

(iii) Real-time Geometry & Motion Reconstruction:

- Deformation Graphs (ED)
- Non-rigid ICP (dense depth + dense colors) instead of SIFT.
- illumination correction

(iv) Fusion 4D:

- way more complex.
- 8 Depth maps (16 IR cameras)
- Key volumes: multiple poses → i.e. allows for topological changes guy wearing coat etc.

(v) Motion 2 Fusion:

- Basically faster Fusion 4D: 200FPS
- ED Graphs
- Key Volumes

* Motion 2 Fusion had 2 novelties: 1) 2 way backward and forward solves it by matching reference to data then (data to reference: $\text{I} \rightarrow \text{J}$)

i.e. solving topology problem

2) Detail layer: they increased resolution of grid at surface (viewing rays)

close to open:
reference to data
open to close:
data to reference.

(vi) Lookin Good: to increase / improve reconstructions using deep architectures

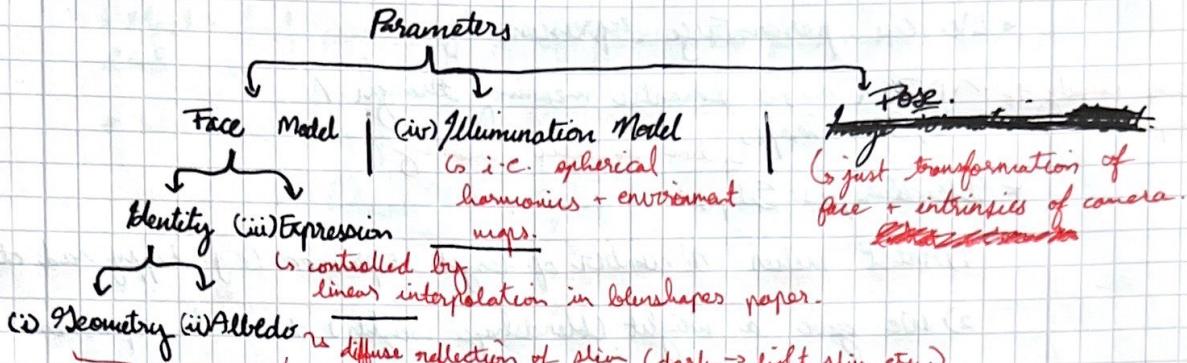
(vii) DeepDeform: to find better correspondences than SIFT for best tracking

The two RGBs are the source query points and target where correspondences need to be found.
The two RGBs are the source query points and target where correspondences need to be found.
use siamese network that shares weights for two RGBs to give you correspondences.

$\rightarrow \Rightarrow$ For non-rigid reconstruction: accumulates drift if tracking is wrong.

Tracking with a)
pre-defined mesh mitigates drift wrt. base mesh.

Lecture 8: Parametric Face Models



(i) Identity: Geometry

* Here by geometry we mean ~~the~~ structure of face (can include expressions)

* Steps ~~the~~: According to face warehouse + morphable model

1) Scan a bunch of faces

2) Fit topologically consistent template ~~non-rigidly~~ to faces

↳ Since all registration from same model we have 1 to 1 ~~one-to-one~~ correspondences

3) Create one massive vector of all positions of vertices per face

4) Find PCA-basis for all the vectors (i.e. all faces)

⇒ $M_{geo}(\alpha) = \text{avg} + E\alpha$ ~~linear parameter~~ ↳ this also means we get average shape ↳ PCA basis that we found.

eigenvectors of the

* The principle components are the n largest eigenvalues of covariance matrix

↳ also means direction of largest variance in our data or ~~line~~ line with minimal dist. to data.

$$\text{Covariance: } \text{cov}(X, X) = \frac{1}{N-1} \sum_i \bar{X}_i \bar{X}_i^T$$

where $\bar{X}_i = X_i - \bar{X}$ is average \bar{X}_i

(ii) Identity: Albedo

* Here we mean light/dark skin

* Steps exactly the same as Geometry but done independently of Geometry.

~~average albedo~~

↳ PCA basis we found

$$M_{alb}(\beta) = \alpha_{alb} + E\beta \quad \text{↳ linear parameter.}$$

(iii) ~~Expression~~ Expression:

* We can parametrize expressions by:

done usually to ~~express~~ 1) PCA as no semantic meaning though /
express data.
2) BlendShapes ~~has semantic meaning~~

* BlendShapes: Steps:

1) Artist makes n number of target expressions (e.g. happy sad etc.)

2) We give a weight (blendshape weights) to each expression
then sum/interpolate all expressions to get final expression

(e.g. 0.5 happy + 0.5 sad.)

(iv) Illumination Model:

Steps

↳ Environment Maps

↳ Spherical Harmonics ~~parametrization of environment maps~~

1) Environment maps place the object a cube/sphere map
of some panoramic view

2) Project lighting unto object (which is in center of sphere/cube)

~~Assumptions~~: → Distant light
→ No scattering

* Spherical harmonics parametrize/reduce dimensionality
of environment maps. (orthogonal basis)

⇒ we can linearly interpolate basis spheres for
different lighting

~~Assumptions~~: → Lambertian Surface:
→ Distant Smooth Light

light
↳ i.e. diffuse
lighting ~~depends on~~
only depends on \vec{n}

Lecture 9: Face Tracking and Reconstruction

* Goal: find parameters \hat{P} of parametric face model to fit RGB 2D image.

* We make energy term: $E(P)$ for energy minimization

$$E(P) = E_{dense}(P) + E_{sparse}(P) + E_{reg}(P)$$

Energy of color and geometry of mesh w.r.t. P Energy Correspondence terms w.r.t. P regularization of P

$$E_{dense}(P) = E_{geom}(P) + E_{col}(P)$$

point to point \hookrightarrow color. point to plane \hookrightarrow point to point

$$E_{geom}(P) = \sum_i \|M_i(P) - D(\pi(M_i(P)))\|_2^2 + [(M_i(P) - D(\pi(M_i(P)))) \cdot n_i]^2$$

summed over all vertices mesh vertex correspondant depth value at projected mesh vertex in camera frame dot product of mesh mesh normal.

$$E_{col}(P) = \sum_i \|M_i(P)_c - I(\pi(M_i(P)))\|_2^2$$

colors of mesh at vertex i (in mesh) color of pixel at projected mesh vertex in camera frame.

$$E_{sparse}(P) = \sum_j \|I(\pi(M_j(P))_m) - C_j\|_2^2$$

2D coordinates of projected mesh in camera frame 2D coordinates of correspondence in 2D image

$$E_{reg}(P) = \sum_k \left(\frac{\alpha_k}{\sigma_{id,k}} \right)^2 + \sum_l \left(\frac{\beta_l}{\sigma_{all,l}} \right)^2 + \sum_m \left(\frac{\gamma_m}{\sigma_{exp,m}} \right)^2$$

α, β, γ are parameter vectors of id, all, exp f centered at zero of. The summands done for each dimension.

$\Rightarrow P^* = \underset{P}{\operatorname{argmin}} E(P)$, solved with non-linear least squares.

\hookrightarrow Gauss Newton or LM solver.

$$\Rightarrow P_{n+1} = P_n - (\mathbf{J}_F^T \mathbf{J}_F)^{-1} \mathbf{J}_F^T \cdot \mathbf{F} \quad \text{Usually done Rough} \rightarrow \cancel{\text{pyramid}} \rightarrow \text{Dense}$$

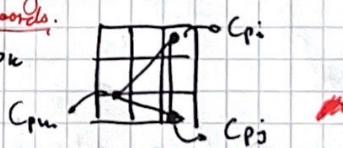
Partial Derivatives of 2D image w.r.t. P :

1) For each color pixel:

color at p_i etc.

$$\text{Pixel} = \alpha C_{pi} + \beta C_{pj} + \gamma C_{pk}$$

Barcentric coords.
Linear interpolation of vertex colors.



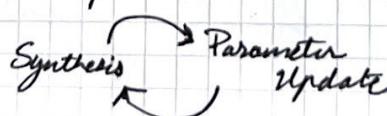
probably the same but not sure :/

2) record triangle v , barcentric coords α, β, γ .

$$3) \Rightarrow \frac{\partial \text{Pixel}(P)}{\partial P_j} = \alpha \frac{\partial C_{pi}(P)}{\partial P_j} + \beta \frac{\partial C_{pj}(P)}{\partial P_j} + \gamma \frac{\partial C_{pk}(P)}{\partial P_j}$$

* This parameter update done:

BRAUNEN



- * Can be done without depth via multiple views
 - ↳ energy jointly optimized for n images

- * Energy optimization dependent on:

- Input Data (RGB vs. RGB-D)
- Face Model + Illumination (coarse vs. dense)
- Run-time requirements, hardware
- Application

- * Two main papers using this:

- Real-time expression transfer for facial reenactment
 - ↳ for two live videos
- Face2Face
 - ↳ expression to RGB video

Lecture 10: Body & Hand Tracking

→ Parameterization for consistent topology of hand + body

→ Parameters of body: $M(\vec{\theta}, \vec{\beta}, \vec{u} | \phi)$

pose of body → hyperparameters.
shape → texture

CAESAR Dataset

→ 125 m 125 F

→ 74 markers.

→ Various ethnicities, age, weight.

→ Same idea as before: 1) average mesh + Non-rigid ICP

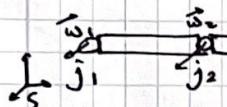
2) PCA of all mesh vectors.

SCAPE: → joint shape + pose learning.

→ Blend shapes used first for poses.



Skeleton Side note:

 P_{j1} : each joint has a transformation matrix related to rotation by Rodrigues formula

$$\Rightarrow p_s = g(\vec{\omega}_1, j_1) g(\vec{\omega}_2, j_2) p_b$$

→ poses are given by $\theta = (\vec{\omega}_1, \dots, \vec{\omega}_n)$ where $\vec{j} = (j_1, \dots, j_n)$

Linear Blend Skinning: For any vertex t_i' on mesh:

vertices are linear combinations of all joints $\in t_i' = \sum_{k=1}^n w_{k,i} M_k(\theta, \vec{j}) t_i$

weight given to each joint

transformation of each joint

Corrective Blend Shapes:

Basically displace each vertex by $P = \begin{bmatrix} \Delta x_1 & \Delta y_1 & \Delta z_1 \\ \vdots & \vdots & \vdots \\ \Delta x_n & \Delta y_n & \Delta z_n \end{bmatrix}$
to simulate muscle compression.)

Non-rigid registration: $E = w_d E_d + w_s E_s + w_m E_m$

$$E_d = \sum_i w_i \text{dist}^2(T_i, V_i, D)$$

Data error.

$$E_s = \sum_i \|T_i - T_j\|_F^2$$

Transformation Smoothness

$$E_m = \sum_i \|T_i v_{ki} - m_i\|^2$$

marker Errors.

Hand Models:

Taylor: Uses subdivision model ~ recursive coarser mesh.

BRUNNEN

to get:

→ C_2 continuity

→ Fast LM convergence

Issues with Model Generation + Fitting

- Expensive Scanning setup
- Neutral Pose similar but not identical
- Ambiguities in fitting \Rightarrow change shape or expression?
- Drift in non-rigid registration