

# INTELLIGENT AGENTS

**Intelligent agent:** Perceives its environment through sensors and acts upon the environment through actuators.

**Percept sequence:** Complete history of its perception (sensor readings). **Agent function:** Maps percept sequence to an action. **Rational agent:** For each possible percept sequence, a rational agent should select an action that is expected to maximize its performance measure, given the prior percept sequence and its built-in knowledge. **Omniscient agent:** Agent knows the actual outcome of its actions (impossible in reality).

**Learning:** Learn from perception, i.e. agent improves their knowledge of the environment over time. **Autonomy:** Agent is considered more autonomous if it is less dependent on prior knowledge and uses newly learned abilities. **Task environment:** Defined by PEAS (performance measure, environment, actuators, sensors). **Properties:** Fully vs. partially observable (agent can/cannot detect complete state of env.). Single vs. multi agent (single/several agents) in env.). Deterministic vs. stochastic (next state is (not) fully determined by its current state and the action of the agent). Episodic vs. sequential (action in one episode does (not) affect later episodes). Discrete vs. continuous (applies to both state & time). Static vs. dynamic (env. only changes based on actions of the agent (not)). Known vs. unknown (agent (not) knows the outcome of its action).

**Agent types:** Simple reflex agent. Selects action only based on current percept. Model-based reflex agent has internal state

that depends on the percept history → can handle partial observability. **Goal-based agent:** Additional has explicit goal → reflected in actions performed by agent; Reaching goal is achieved by search and planning [ex: "reach final destination"]

Utility-based ag. Additionally tries to maximize utility [ex: "reach final dest. in 25 min."]. All can be turned into Learning agents. Tree search: actions not reversible → finite

UNINFORMED SEARCH: Acyclic transition graph → finite search tree.

Properties: No additional information besides problem statement. Well defined problems: initial state, actions, transition model, goal test,

path cost. Tree search: Current node + frontier list → can get stuck in loops. Graph search: Additionally has "explored set", so a node cannot be visited twice.

Notation branching factor  $b$  (max number of successors of any node), depth  $d$  of shallowest goal node, Max: length  $m$  of any path in the state space, cost  $C^*$  of optimal solution,  $E$  minimum step cost

Criterion	BFS	Dijkstra	DFS	IDS
Complete?	Yes <sup>1,2</sup>	No <sup>*</sup>	Yes, if $I=d$	Yes <sup>1</sup>
Optimal?	Yes	No	No, if $I=d$	Yes <sup>3</sup>
Time	$O(b^d)$	$O(b^{1+(C/E)})$	$O(b^m)$	$O(b^d)$
Space	$O(b^d)$	$O(b^{1+(C/E)})$	$O(bm)$ , $O(b)$	$O(bd)$

1: complete if  $b$  is finite (and  $d$ ) 2: complete if step costs  $\geq E$  for pos.  $E$ . 3: optimal if all step costs are identical

FIFO Breadth-first search (BFS): Expands shallowest nodes first. Goal test performed for children of current node. Uniform-cost search (Dijkstra) Order frontier as priority queue, ordered by  $g \rightarrow$  Expands node with lowest path cost  $g(n)$ . Goal test performed when node selected for expansion (+ BFS). Depth-first search (DFS) Expands deepest unexpanded node first. \*Not complete if recursively implemented; Complete if repeated states are avoided and state space is finite. Does not terminate in infinite state spaces.

LIFO Depth-limited search: I: depth limit Iterative deepening search (IDS) Calls DFS with increasing I. Bidirectional search

Run two searches: One from initial state and one from goal. Motivation:  $b^{d/I} + b^{d/I} < b^d$ . Difficult when goal state is unknown or there are multiple goals or the actions are not reversible.

function Breadth-First/Depth-First (problem) returns a solution, or failure

node ← a node with state = problem, initial=state. If problem.goal-test(node.state) then return Solution(node). frontier ← a FIFO/LIFO queue with node as the only element. explored ← an empty set

loop do  
if Empty(frontier) then return failure  
node ← Pop(frontier) chooses the shallowest/deepest node in frontier. I add node.state to explored

for each action in problem.actions(node.state) do  
child ← Child-node(problem, node, action)  
if child.state is neither in explored nor frontier then

if problem.goal-test(child.state) then return Solution(child). frontier ← Insert(child, frontier)

function Depth-Limited-S. (problem, limit) returns a solution, or failure/cutoff

return Recursive-DLS(Make-Node(problem, initial-state), problem, limit)

function Recursive-DLS(node, problem, limit) returns a solution, or failure/cutoff

if problem.goal-test(node.state) then return Solution(node). else if limit = 0 then return cutoff

cutoff.occurred ← false  
for each action in problem.actions(node.state) do

child ← Child-node(problem, node, action)  
result ← Recursive-DLS(child, problem, limit - 1)

if result = cutoff then cutoff.occurred ← true  
else if result ≠ failure then return result

if cutoff.occurred then return cutoff else return failure

INFORMED SEARCH

Properties Know whether a state is "better" than another to reach a goal by using an evaluation function:  $f(n) = g(n)/h(n)$

g: sum of path cost, h: estimated cost current node → goal,  $h(n) \geq 0$

$h(goal) = 0$ ; all informed search alg. are based on Dijkstra's alg.

**Greedy best-first search**: Expands node closest to goal by using heuristic func  $f(n) = h(n)$ . Completeness: Yes, w/ graph search, else not. Not optimal. Time  $O(b^m)$  (heuristic misleading, solution found last). Space  $O(b^m)$ .  $A^*$  search:  $f(n) = g(n) + h(n)$

→ never overestimate the cost. Admissible:  $h(n) \leq g(n, goal)$  needed for optimality of tree search. Consistent:  $h(n) = g(n, n') + h(n')$  → needed for optimality of graph search (stronger). Completeness: Yes, if costs > 0. Optimal if  $h(n)$  is valid. Time  $O(b^d)$ . Space  $O(b^d)$ .  $E = \frac{h}{h} = \frac{h}{h}$ ,  $h/h$ : actual/estimated cost from root to goal. Heuristic function Effective branching factor  $b^*$

avg. number of children in search tree → lower  $b^*$ , better heuristic. Domination:  $h_2$  dominates  $h_1$  if  $h_2(n) \geq h_1(n)$  for all nodes  $n \in A^*$  using  $h_2$  will never expand more nodes than with  $h_1$ .

Obtain heuristics e.g. from related problems or pattern databases ground function Uniform-Cost/Greedy-Best-First/A\* (problem) returns a solution, or failure

node ← a node with state = problem, initial=state. frontier ← a priority queue by Evaluation-Function with node as the only element. explored ← an empty set

loop do

if Empty(frontier) then return failure

node ← Pop(frontier) chooses the node in frontier with lowest Evaluation-Function if problem.goal-test(node.state) then return Solution(node). I add node.state to explored

for each action in problem.actions(node.state) do

child ← Child-node(problem, node, action)

if child.state is neither in explored nor frontier then

← frontier ← Insert(child, frontier)

else if child.state is in frontier with higher Evaluation-Function then

← replace that frontier node with child

Graph search w/ Tree search modifications

Evaluation-Function<sup>1</sup>

$P_i, P_j, P_k, P_l, P_m, P_n, P_o, P_q, P_r, P_s, P_t, P_u, P_v, P_w, P_x, P_y, P_z$

$Subst(O, p) = Subst(G, q), e.g. O = \exists x \exists y / E \exists x \exists y / E$

FOL Horn Clauses Like in PL + universally quantified vars ("V" is typically omitted)

Backward Chaining algorithm Start from goal, search rules to find known facts that support the proof. Needs Horn clauses (then complete).

goals:  $\{goal, \dots\}$  applied. Initialize with sentence to prove.

$\emptyset \leftarrow \{goal, \dots\}$  Current substitution. Initialize with empty set  $\emptyset$ .

$Q \leftarrow Subst(\emptyset, goal) \leftarrow q^*$ :

Using rule i: sentence<sub>i</sub>, ..., sentence<sub>n</sub> → goal Standardize apart

$\emptyset \leftarrow \{x_i / object\}$  Current unification

new\_goals ← {sentence<sub>1</sub>, ..., sentence<sub>n</sub>, old\_goal, ...}

if goals empty → successful proof

Conversion to CNF 1) Eliminate implications ( $A \Rightarrow B = \neg A \vee B$ )

2) Move → inwards 3) Standardize vars ( $x \rightarrow x_i$ )

4) Skolemization ( $\forall t \exists p \text{ Parent}(p, c) \rightarrow \exists \text{ Parent}(t, c)$ )

5) Drop universal quantifiers & 6) Distribute V over ∧

Resolution  $\{P_1, \dots, P_n, Q_1, \dots, Q_m\} \vdash_{\text{Unify}(P_1, Q_1)} \dots \vdash_{\text{Unify}(P_n, Q_n)}$  with  $\emptyset = \text{Unify}(P_1, Q_1)$

rule  $Subst(\emptyset, P_1, \dots, P_n, Q_1, \dots, Q_m)$

Resolution algorithm Like in PL but important to standardize apart. In FOL, resolution can always prove that a sentence is unsatisfiable. Attempting to prove a satisfiable first-order formula as unsatisfiable may result in a non-terminating computation (unlike in PL).

BAYESIAN NETWORKS static environment, without actions

Bayesian network Directed acyclic graph, where each node

corresponds to a random variable, arrows between nodes start at parents. Each node  $N$  has a conditional probability distribution  $P(X_i | \text{Parents}(X_i))$ . Bayes rule  $P(X|Y) = \frac{P(Y|X)P(X)}{P(Y)}$

Independence  $P(X, Y) = P(X)P(Y) \Leftrightarrow$  Variables  $X$  and  $Y$  share no common ancestry. Conditional independence  $X$  conditionally independent of  $Y$  given  $E$ :  $P(X|Y|E) = P(X|E)P(Y|E) \Leftrightarrow$

$P(X|Y, E) = P(X|E) \Leftrightarrow X$  and  $Y$  are independent given  $E$  iff every path  $U$  from  $X$  to  $Y$  is blocked, i.e. for every path  $U$  after all manipulations there is no undirected path between  $X$  and  $Y$ .

Constraint Satisfaction Problems CSP ( $X, D, C$ )  $X = \{x_1, \dots, x_n\}$  variables,  $D = \{D_1, \dots, D_n\}$  domain of vars,  $C = \{C_1, \dots, C_m\}$  constraints,  $D_i = \{v_1, \dots, v_n\}$  allowable values. Constrained types Unary (1 var involved), Binary (2 vars), Higher-Order ( $> 2$  vars). Convert n-ary constraint to binary constraint 1. Replace  $l$ -ary constraint with var  $Z$  with domain of  $l$ -tuple and relation  $R(X_1, X_2, \dots, X_l)$ :  $\text{dom}(Z) = \{(z_1, \dots, z_l) | R(z_1, \dots, z_l), z_i \in \text{dom}(X_i), z_1, \dots, z_l \in \text{dom}(X_l)\}$  2. Introduce new binary constraints to match values of  $Z$  with values of neighboring vars:  $C = \{X_1, X_2, \dots, X_l\}$

...  $\langle (X_1, X_2), l^h(Z) = X_3 \rangle\}$  3. New CSP:  $X = X \cup Z$ ,  $D = D \cup \text{dom}(Z)$ ,  $C'$  [Hidden transformation] Backtracking heuristics Minimum Remaining Value (MRV) (Variable Selection) Assign Var with fewest possible values first. Degree heuristic (Variable Selection) Select variable involved in the largest number of constraints on other unassigned vars.

Least Constraining Value (Value Selection) Choose value that rules out fewest choices for neighboring values in the constraint graph.

Backtracking Search DFS for CSPs with single variable assignment

in each step 1. Select Var to assign 2. Set value of domain

3. Check consistency (inference) → add new assignment or backtrack and try new assignment if you find inconsistency. Inference algorithms

Forward Checking (After each assignment): inconsistent values of neighboring vars are removed. Arc consistency algorithm (AC-3) (also as pre-processing): inconsistent values of all vars are removed.

function AC-3 (csp, queue) returns failure or the reduced csp otherwise

inputs: csp: a binary CSP, queue: a queue of arcs ( $(X_i, X_j)$ ) while queue is not empty do

$(X_i, X_j) \leftarrow \text{Remove-Inconsistent-Values}(X_i, X_j)$  then

if size of  $\text{Domain}(X_i) = 0$  then return failure

for each  $X_k$  in  $\text{Neighbors}(X_i) \setminus \{X_j\}$  do

add  $(X_k, X_i)$  to queue

return csp

function Remove-Inconsistent-Values ( $X_i, X_j$ ) returns true iff succeeds

removed ← false

for each  $x \in \text{Domain}(X_i)$  do

if no  $y \in \text{Domain}(X_j)$  allows  $(x, y)$  to satisfy the constraint of  $(X_i, X_j)$  then

then delete  $x$  from  $\text{Domain}(X_i)$ ; removed ← true

return removed

Initialization of variable "queue" in AC-3: 1) After each assignment: After assigning Var  $X_i$ , add the arcs  $(X_i, X_j)$  to queue ( $X_i$ : all unassigned neighbors of  $X_i$ ). 2) As pre-processing: add all arcs of the CSP to queue.

arc consistency 1) Variable:  $X_i$  is arc-consistent with  $X_j$ , if for every value in  $D_i$ , there exists a value in  $D_j$  satisfying the binary constraint of the arc  $(X_i, X_j)$ . 2) CSP: A constraint graph is arc-consistent if every Var is arc-cons. with every other Var. Independent subproblems

Completely disconnected in constraint graph. Tree-structured CSPs (if graph has no loops): Can be solved in  $O(n^d)$  time (vs.  $O(d^n)$  in general)

→ Use Conditioning and Tree decomposition for transformation into tree.

LOGICAL AGENTS/PROPOSITIONAL LOGIC

Satisfaction If sentence  $\alpha$  is true in model  $m \rightarrow m$  satisfies  $\alpha$ .

KB satisfiable  $\Leftrightarrow M(KB) \neq \emptyset$  [M(a): set of all models satisfying  $\alpha$ ]

α Entailment Truth of one sentence requires the truth of the other sentence:  $\alpha \models \beta$  if  $\alpha$  entails  $\beta$ . Or:  $\alpha \models \beta \Leftrightarrow M(\alpha) \subseteq M(\beta)$

$(x = 0) \models (x = 0)$ . Inference by enumeration Enumerate all models. KB =  $\alpha \models \alpha$  is true in all models in which KB is true. Theorem

proving Apply inference rules Logical equivalences  $\alpha \equiv \beta$  iff  $\alpha \models \beta$  and  $\beta \models \alpha$ .  $\alpha \models \beta \equiv (\neg \beta \rightarrow \neg \alpha)$  Validity/Tautology Sentence is valid iff true for all models.  $M(\alpha \wedge \neg \alpha) = M(\text{true}) = \text{All Possible Models}$

$\alpha$  valid  $\Leftrightarrow \neg \alpha$  unsatisfiable Satisfiability Sentence is satisfiable if true for some models. Unsatisfiable:  $M(\alpha \wedge \neg \alpha) = M(\text{false})$   $\models \neg \alpha$  If  $\alpha \models \beta$  then  $\neg \beta \models \neg \alpha$

rules Modus Ponens  $\frac{\alpha \models \beta}{\alpha \wedge \beta \models \beta}$  And-Elimination  $\frac{\alpha \wedge \beta}{\alpha \models \beta}$  Conjunction  $\frac{\alpha \models \beta \quad \alpha \models \gamma}{\alpha \models \beta \wedge \gamma}$  Normal form (CNF)  $(A \vee B \vee C) \wedge (\neg A \vee \neg C)$  Resolution algorithm

To prove  $KB \models \alpha$ , we show  $KB \models \neg \alpha \rightarrow \perp$  is unsatisfiable. Apply resolution rule  $\frac{A \vee B \quad C \vee D}{A \vee C \quad B \vee D}$  to CNF until: 1) no new clause produced  $\rightarrow KB \models \alpha$

2) Clauses resolve to yield empty clause  $\perp \rightarrow KB \models \alpha$  proof by contradiction

Horn clauses

$(P_1 \wedge \dots \wedge P_n) \rightarrow Q$ ,  $Q$  and  $P_i$ : propositional symbols. KB of Horn clauses

$(P_1 \wedge \dots \wedge P_n) \rightarrow Q$   $\models P_1, \dots, P_n \rightarrow Q$   $\models P_1, \dots, P_n \rightarrow Q$

INFORMATION RETRIEVAL

Properties Know whether a state is "better" than another to reach a goal by using an evaluation function:  $f(n) = g(n)/h(n)$

g: sum of path cost, h: estimated cost current node → goal,  $h(n) \geq 0$

$h(goal) = 0$ ; all informed search alg. are based on Dijkstra's alg.

only requires Modus Ponens for inference → Complete for Horn clauses

AND-OR graph  $A \wedge B \Rightarrow C, C \Rightarrow D \Rightarrow A \wedge B \Rightarrow D$  Forward+Backward Chaining Linear time. Data/goal driven. May do irrelevant work.

FIRST-ORDER LOGIC (FOL)

Quantifiers  $\forall$  goes with  $\wedge$  and  $\exists$  with  $\wedge$ . Logic translation

$\forall x P(x) \rightarrow Q(x)$  All Ps are Qs.  $\exists x P(x) \rightarrow \neg Q(x)$  No Ps are Qs.

$\exists x P(x) \wedge Q(x)$  Some Ps are Qs.  $\forall x P(x) \wedge \neg Q(x)$  Some Ps aren't Qs.

Syntax AtomicSentence ::= Predicate(PredicateTerm, ... ) Term = Term;

Term ::= Function(Term, ... ) | Constant | Variable De Morgan rules

quantifiers  $\forall x P \equiv \neg \exists x \neg P$ ,  $\exists x P \equiv \neg \forall x \neg P$ ,  $\neg \exists x P \equiv \forall x \neg P$

predicates represent relations

INFERENCE IN FOL

Removing Quantifiers  $\forall$ : Infer any sentence from substituting a bound term for the variable

$\forall x f(x, g(x, h(x))) \rightarrow f(g(h), g(h))$  (several times applicable)

$\exists x f(x, g(x, h(x))) \rightarrow f(g(h), g(h))$  (once applicable)

Proportionalization Instantiate universal sentence in all possible ways Generalized Modus Ponens For atomic sentences  $P_i, P_j, P_k, P_l, P_m, P_n, P_o, P_q, P_r, P_s, P_u, P_v, P_w, P_x, P_y, P_z$

such that  $Subst(O, P_i) = Subst(O, P_j)$ : Horn clauses

$P_i, P_j, P_k, P_l, P_m, P_n, P_o, P_q, P_r, P_s, P_u, P_v, P_w, P_x, P_y, P_z$

$Subst(O, P_i) = Subst(G, q), e.g. \exists x \exists y / E \exists x \exists y / E$

FOL Horn Clauses Like in PL + universally quantified vars ("V" is typically omitted)

Backward Chaining algorithm Start from goal, search rules to find known facts that support the proof. Needs Horn clauses (then complete).

goals:  $\{goal, \dots\}$  applied. Initialize with sentence to prove.

$\emptyset \leftarrow \{goal, \dots\}$  Current substitution. Initialize with empty set  $\emptyset$ .

$Q \leftarrow Subst(\emptyset, goal) \leftarrow q^*$ :

Using rule i: sentence<sub>i</sub>, ..., sentence<sub>n</sub> → goal Standardize apart

$\emptyset \leftarrow \{x_i / object\}$  Current unification

new\_goals ← {sentence<sub>1</sub>, ..., sentence<sub>n</sub>, old\_goal, ...}

if goals empty → successful proof

Conversion to CNF 1) Eliminate implications ( $A \Rightarrow B = \neg A \vee B$ )

2) Move → inwards 3) Standardize vars ( $x \rightarrow x_i$ )

4) Skolemization ( $\forall t \exists p \text{ Parent}(p, c) \rightarrow \exists p \text{ Parent}(t, c)$ )

5) Drop universal quantifiers & 6) Distribute V over ∧

Resolution rule  $\{P_1, \dots, P_n, Q_1, \dots, Q_m\} \vdash_{\text{Unify}(P_1, Q_1)} \dots \vdash_{\text{Unify}(P_n, Q_n)}$

Resolution algorithm Like in PL but important to standardize apart. In FOL, resolution can always prove that a sentence is unsatisfiable. Attempting to prove a satisfiable first-order formula as unsatisfiable may result in a non-terminating computation (unlike in PL).

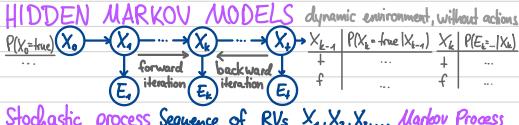
Bayesian Networks static environment, without actions

Bayesian network Directed acyclic graph, where each node

corresponds to a random variable, arrows between nodes start at parents. Each node  $N$  has a conditional probability distribution

$P(X_i | \text{Parents}(X_i))$  Markov blanket: A node in the Bayesian network is independent of all other nodes given its parents, children and children's parents.

Inference by Enumeration  $P(A|d, e) = P(A|d, e)$



**Stochastic process** Sequence of RVs  $X_1, X_2, X_3, \dots$  **Markov Process**  
**Stochastic process** that has the **Markov property** **Markov property**  
 $P(X_n | X_{n-1}, X_{n-2}, \dots, X_0) = P(X_n | X_{n-1})$  **Stationary process** Stochastic process whose joint prob distribution does not change when shifted in time. For Markov processes:  $\forall i: P(X_n | X_{n-i}) = P(X_{n+i} | X_{n+i})$   
**Stationary Markov Chain** Discrete stat. process with Markov prop.  $P(X_n = x_i) = \sum_{j=1}^k P(X_1 = x_j) \cdot P(X_{n-1} = x_j)$ . Matrix notation:  $P_n = T \cdot P_{n-1}$ ,  $(P_n)_{ij} = P(X_n = x_i, X_{n-1} = x_j)$  Convert  $m^{th}$ -order MCs to first-order: Introduce new states corresponding to the  $n$ -th previous states:  $P(X_n = x_i | X_{n-1} = x_1, X_{n-2} = x_2, \dots) = P(X_n = x_2 | X_{n-1} = x_3)$  with  $x_2 = (x_1, x_2), x_3 = (x_2, x_3)$  **HMM** = Markov Chain + sensor model. State cannot be measured directly  $\rightarrow$  inferred from sensor values (which only depend on current state):  $P(E_t | X_t)$   
**Basic formulas**  $(P_i)_k = P(X_k = x_i)$ ,  $(\bar{P}_i)_k = P(E_k = e_i)$

$$\text{dim}(T) = |\mathcal{S}| \times |\mathcal{S}|^m \quad T_{i,j} = P(X_i = x_i | X_{i-1} = x_j) = \begin{cases} P(x_i | x_{i-1}) & P(x_i | x_{i-1}) \\ P(x_i | x_{i-1}) & P(x_i | x_{i-1}) \end{cases}$$

$$P_L = T \cdot P_{L-1}, \quad \bar{P}_L = H \cdot P_L, \quad H_{i,j} = P(E_k = e_i | X_k = x_j)$$

Joint prob. dist.:  $P(X_{t+1}, E_{t+1}) = (\prod_{i=1}^m P(E_i | X_i)) \cdot P(X_{t+1}) \cdot P(X_t)$

**Filtering** Recursive state estimation algorithm. Time:  $O(t)$ . Space:  $O(1)$

$$f_{1:t} = P(X_1 | e_{1:t}) = \alpha \cdot P(e_1 | X_1) \cdot \sum_{X_1} P(X_1 | X_{1:t}) \cdot P(X_{1:t} | e_{1:t-1})$$

$$= \alpha \cdot O_1 \cdot T \cdot f_{1:t-1}$$

$$(f_i)_{1:t} = P(X_1 = x_i | e_{1:t}), \text{ initialization: } (f_i)_{1,0} = (f_i)_{1,0} = P(X_1 = x_i)$$

$$(O_{ij})_{1:t} = \begin{cases} P(e_j | X_1 = x_i), \text{ if } i=j \\ 0, \text{ else} \end{cases} = \begin{cases} P(e_j | X_1) & 0 \\ 0 & P(e_j | X_1) \end{cases}$$

**Prediction** Like filtering without adding evidence (inaccurate for high  $t$ )

$$P(X_{t+k} | e_{1:t}) = \sum_{X_{t+k}} P(e_{t+k} | X_{t+k}) \cdot P(X_{t+k} | e_{1:t}) = T^{k+1} f_{1:t}$$

**Smoothing** Computing the distribution over past states given evidence up to the present.  $P(X_k | e_{1:t})$  for  $0 \leq k \leq t$ .

$$P(X_k | e_{1:t}) = \alpha \cdot P(X_k | e_{1:t}) \cdot P(e_{t+k+1} | X_k) = \alpha \cdot f_{1:k} \cdot b_{k+1:t}$$

**Backward recursion:** Initialization with  $b_{k+1:t} = (1, \dots, 1)^T$

$$b_{k+1:t} = P(e_{t+k+1} | X_k) = \sum_{X_{t+k}} P(e_{t+k+1} | X_{t+k}) \cdot P(e_{t+k+1} | X_{t+k}) \cdot P(X_{t+k} | X_k)$$

$$= T^T \cdot O_{k+1:t} \cdot b_{k+2:t} \quad (b_i)_{k+1:t} = P(e_{t+k+1} | X_k = x_i) \quad \text{Begin: } k=t-1$$

For whole sequence time complexity  $O(t^2)$ , can make it  $O(t)$  using dynamic programming ( $\rightarrow$  forward-backward algorithm)

**Viterbi algorithm** Determine the most likely sequence of states. Time:  $O(t)$ . Space:  $O(t)$ .

$$\text{Initialize: } u_i(X_i) = f_{1,i}$$

$$u_{t+1}(X_{t+1}) = \max_{x_1, \dots, x_t} P(x_1, \dots, x_t, X_{t+1} | e_{1:t+1}) = u_t(X_t)$$

$$= \alpha \cdot P(e_{t+1} | X_{t+1}) \cdot \max_{x_1, \dots, x_t} P(X_{t+1} | x_t) \cdot \max_{x_1, \dots, x_t} P(x_t | x_{t-1})$$

Matrix notation:

$$u_{t+1}(X_{t+1}) = \alpha \cdot O_{t+1} \cdot \begin{bmatrix} \max_{x_1} \{ T_{11} \cdot u_1(x_1), \dots, T_{1n} \cdot u_1(x_1) \} \\ \vdots \\ \max_{x_1} \{ T_{n1} \cdot u_n(x_1), \dots, T_{nn} \cdot u_n(x_1) \} \end{bmatrix}$$

$$u_{t+1}(X_{t+1}) = \alpha \cdot O_{t+1} \cdot \begin{bmatrix} \max_{x_1} \{ T_{11} \cdot u_1(x_1), (T_{12} \cdot u_1(x_1)) \} \\ \vdots \\ \max_{x_1} \{ T_{21} \cdot u_1(x_1), (T_{22} \cdot u_1(x_1)) \} \end{bmatrix}$$

Pick highest final state and backtrace! ( $\alpha$  can be omitted)

$$X = \text{true} \quad \begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \end{array} \quad \begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \end{array} \quad \begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \end{array} \quad \begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \end{array}$$

$$X = \text{false} \quad \begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \end{array} \quad \begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \end{array} \quad \begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \end{array} \quad \begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \end{array}$$

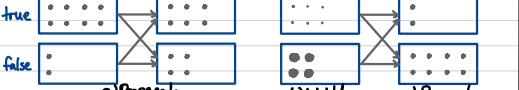
Summarize Evidence Variables e.g.  $A, B \in \mathbb{S} \cdot \mathbb{F}^k$  to  $\text{for } k \geq 0$

$$E = \{(a, b), (a, \neg b), (\neg a, b), (\neg a, \neg b)\}$$

**Steady state condition:**  $f_{t+k} = f_{t+k-1} = f_{t+k-2} = \dots = f_{t+1}$

$$f_{t+k} = \alpha \cdot O_k \cdot T \cdot f_{t+k-1}$$
 (Filtering formula)  $\rightarrow$  solution: Eigenvector  $v \rightarrow f_{t+k} = \tilde{\alpha} \cdot v$ 

**Particle Filtering** (Monte Carlo for HMM): Population of randomly initialized particles track high-likelihood regions of state-space.



a) Propagation through the transition model b) Evidence observed, weight each sample by its likelihood for the observation c) Generate new set of samples by weighted random selection from current set

**RATIONAL / SIMPLE DECISIONS** actions in static envs

$P(\text{Result}(a) = s' | e) = P(s' | a, e) = \sum_s P(s) \cdot P(s' | s, a, e)$ , evidence  $e$ , action  $a$

$\text{EU}(a|e) = \sum_s P(\text{Result}(a) = s' | a, e) \cdot U(s')$ ,  $\text{MEU}(a|e) = \max_a \text{EU}(a|e)$

$\text{EU}(x) = \sum_s \text{EU}(x|s) = \sum_s P(x|s) \cdot U(s)$

$\text{EU}(x|y) = P(x|y) \cdot U(x|y) = \sum_s P(x|y, s) \cdot U(x|y, s)$

**Decision networks** Decision, RV,  $\bigcup$  Bayesian network with decision and utility nodes Partial ordering 1) Identify the first decision  $D_1$  and all vars  $X_0$  to make that decision 2) Identify next decision  $D_2$  and the vars  $X_1$  that are revealed between  $D_1$  and  $D_2$  3) Place unrevealed vars  $X_n$  at the end of the ordering  $\rightarrow X_0 < D_1 < X_1 < D_2 < \dots < X_n$ ,  $X_i$  may be a set  $\{ \dots \}$  of vars

Fundamental links Removal would not change partial ordering

No forcing assumption: All past decisions and revealed vars are available at the current decision (only need to draw non-fund. links)

Decision trees Nodes as decision networks. Ell of a decision:  $\frac{9}{500} \cdot Q \cdot \frac{9}{50} \cdot \frac{9}{50} \cdot \frac{9}{50} \cdot Q \cdot \frac{9}{50} \cdot \frac{9}{50} \cdot Q$

weighted summation of all branches from decision to all reachable leaves from the decision,  $\rightarrow$  exponential complexity with increasing num. of decisions

Value of information Guides an agent to choose what information to acquire  $\text{VOI}_e(E_i) = [\sum_j P(E_j = e_j | e_i) \cdot \text{MEU}(e_j | e_i, E_j = e_j)] - \text{MEU}(e_i)$

= expected value of best action given the information at no charge - expected value of best action without information Optimal policies

$\pi^*(D_i | X_{1:i-1}, d_{1:i-1}) = \arg \max_{\pi_i} \text{EU}(D_i | X_{1:i-1}, d_{1:i-1})$   $x_i$ : chance nodes

$d_i$ : decision nodes Joint probability  $P(X_{1:n}, d_{1:n}) = P(X_{1:n}, d_{1:n})$

$\pi^*(D_i | X_{1:i-1}, d_{1:i-1})$  Decision maximizing the EU for each decision var.

$\text{MEU}(d_{1:n}) = \max_{d_1} \sum_{x_1} \max_{d_2} \sum_{x_2} \dots \max_{x_n} \sum_{d_n} \text{EU}(X_{1:n}, d_{1:n})$

$\text{EU}(d_n) = \max_{d_n} \sum_{x_n} P(x_n | x_{n-1}, d_{1:n-1}) \cdot \text{EU}(x_n, d_n)$

$= \sum_{x_n} P(x_n | d_n) \cdot \max_{d_{n-1}} \sum_{x_{n-1}} P(x_{n-1} | x_n, d_{1:n-1}) \cdot \text{EU}(x_{n-1}, d_{1:n-1})$

$= \sum_{x_n} P(x_n | d_n) \cdot \max_{d_{n-1}} \text{EU}(d_{n-1} | x_n, d_n)$

1) Calc  $\text{EU}(d_i | X_{1:i-1}, d_{1:i-1})$  for all possible instantiations of vars  $X_{1:i-1}, d_{1:i-1}$

(begin with last decision  $d_n$ ) 2) Derive policy  $\pi^*(D_i | X_{1:i-1}, d_{1:i-1})$  for all possible instantiations of vars  $X_{1:i-1}, d_{1:i-1}$  3) Continue with step 1) with  $d_{i-1}$  and calc EU's by using the obtained policies ( $\rightarrow$  replace  $D_i$  with  $\pi^*(D_i | X_{1:i-1}, d_{1:i-1})$  or directly use the max. EU's.

**RATIONAL DECISIONS OVER TIME** (Complex Decisions) actions in dyn. envs

**Markov Decision Process (MDP)** Sequential decision problem for a fully observable, stochastic env with Markovian transition model and additive rewards (MDP = markov chain + actions + rewards) Definition A MDP

is a 5-tuple  $(S, A, P, R, \gamma)$ , where  $S$  is a finite set of states,  $A$  is

a finite set of actions (or  $A(s)$  is the finite set of actions available from state  $s$ ), transition model  $P(s'|s, a) = P(s_{t+1} = s' | s_t = s, a_t = a)$ ,

reward  $R(s'|s, a)$ , discount factor  $\gamma \in [0, 1]$  (represents the difference in importance between future and present rewards).

Aim Find optimal policy, i.e. the best action for every possible state  $s \rightarrow$  the optimal policy maximizes the expected utility (EU)

Utility/value of a state:  $U(s) = \text{expected (discounted) sum of rewards (until termination)}$  if you start in state  $s$  assuming optimal actions.

$Q$ -function  $Q(s, a) = (\text{Expected}) \text{ Utility when starting in a given state with a given action. Bellmann Principle of Optimality}$

An optimal policy has the property that whatever the initial state and initial decision are, the remaining decisions must constitute an optimal policy with regard to the state resulting from the first decision. Bellman equation

Expected sum of rewards = current reward +  $\gamma \times$  expected sum of rewards after taking best action

$U(s) = R(s) + \gamma \cdot \max_{a \in A(s)} \sum_{s'} P(s' | s, a) \cdot U(s')$  (simplified)

$= \max_{a \in A(s)} \sum_{s'} P(s' | s, a) \cdot [R(s, a, s') + \gamma \cdot U(s')]$

Value Iteration Algorithm 1) Start with arbitrary utility values (except for terminal states, there  $U(s) = R(s)$ )  $\rightarrow$  initial utilities normally  $U(s) = 0$

2) Repeat for every  $s$  simultaneously until "no change" (Bellman Update):

$U^{t+1}(s) = R(s) + \gamma \cdot \max_{a \in A(s)} \sum_{s'} P(s' | s, a) \cdot U^t(s')$  (simplified)  $\rightarrow$  calculate  $U(s)$  separately for all actions  $a \in A(s)$

$= \max_{a \in A(s)} \sum_{s'} P(s' | s, a) \cdot [R(s, a, s') + \gamma \cdot U^t(s')]$

Optimal policy calculate this for all  $a \in A(s)$

$\pi^*(s) = \arg \max_{a \in A(s)} \sum_{s'} P(s' | s, a) \cdot U(s')$  (simplified)

$= \arg \max_{a \in A(s)} \sum_{s'} P(s' | s, a) \cdot [R(s, a, s') + \gamma \cdot U(s')]$

Policy iteration Possible to get an optimal policy even when the utility function estimate is inaccurate ( $\rightarrow$  normally less computation needed)

1) Start with an arbitrary initial policy  $\pi_0(s)$  for every state  $s$ .

2) Policy evaluation: Solve the linear system to compute  $U_i(s)$  for all  $s$

$U_i(s) = R(s) + \gamma \cdot \sum_{s'} P(s' | s, \pi_i(s)) \cdot U_i(s')$  (simplified)

$= \sum_{s'} P(s' | s, \pi_i(s)) \cdot [R(s, \pi_i(s), s') + \gamma \cdot U_i(s')]$

3) Policy improvement for each state  $s$ : calculate this for all  $a \in A(s)$

$\pi_{i+1}(s) = \arg \max_{a \in A(s)} \sum_{s'} P(s' | s, a) \cdot U_i(s')$  (simplified)

$= \arg \max_{a \in A(s)} \sum_{s'} P(s' | s, a) \cdot [R(s, a, s') + \gamma \cdot U_i(s')]$

The algorithm terminates when the policy improvement yields no change in the utilities.

Logical equivalences

$\alpha \wedge \beta \equiv \beta \wedge \alpha$

$\alpha \vee \beta \equiv \beta \vee \alpha$

$(\alpha \wedge \beta) \wedge \gamma \equiv \alpha \wedge (\beta \wedge \gamma)$

$(\alpha \vee \beta) \vee \gamma \equiv \alpha \vee (\beta \vee \gamma)$

$\neg(\neg \alpha) \equiv \alpha$

$\alpha \rightarrow \beta \equiv \neg \beta \rightarrow \neg \alpha$  Better use quantifiers without negation or only negate first quantifier!

$\alpha \Rightarrow \beta \equiv \neg \alpha \vee \beta$

$\neg(\alpha \wedge \beta) \equiv \neg \alpha \wedge \neg \beta$

$\alpha \wedge (\beta \vee \gamma) \equiv (\alpha \wedge \beta) \vee (\alpha \wedge \gamma)$

$\alpha \vee (\beta \wedge \gamma) \equiv (\alpha \vee \beta) \wedge (\alpha \vee \gamma)$

Performance measure vs. Utility function Performance measure is a specification from the designer or programmer (internal to the agent) to specify what the agent should do. If an agent always acts to maximize or achieve the performance measure, this is a rational agent. Not all agents do have a utility function, e.g. roller agents do not. The utility func is used internally by the robot itself to evaluate the best course of action to optimally achieve/maximize the performance measure(s), given its perceived state.

LEARNING An agent is learning if it improves its performance on future tasks after making observations about the world. Types of Learning Supervised Agent observes example input-output pairs and learns a function that maps from input to output. Unsupervised Agent learns patterns in data without feedback (labels). Reinforcement Agent learns from a series of events and receives either reward or punishment for its actions. Occam's razor for supervised learning: Choose the simplest consistent hypothesis Decision trees Input  $\in \mathbb{R}^n \rightarrow$  Output  $\in \mathbb{E}$  tree file? DT structure An internal node represents a test of a property. Edges are annotated with the possible test values. Each leaf node has the Boolean value which should be returned Building DTs For  $n$  boolean attributes:  $2^n$  possible DTs  $\rightarrow$  use greedy divide & conquer heuristic to test most important attribute. First Quantifying importance of attributes Entropy of boolean RV  $X$  that is true with prob.  $q: H(X) = B(q) = -q \cdot \log_2(q) - (1-q) \cdot \log_2(1-q)$  Information gain of an attribute  $A$  with  $d$  values and binary classification Gain( $A$ ) =  $B\left(\frac{P}{P+d}\right) - \sum_{i=1}^d \frac{P_i \cdot n_i}{P+n} B\left(\frac{P_i}{P+n}\right)$

$p$  and  $n$ : number of pos./neg. samples

$p_i$  and  $n_i$ : number of pos./neg. samples of  $k^{\text{th}}$  value

$B\left(\frac{p}{p+n}\right) = 0.3483$ ,  $B\left(\frac{n}{p}\right) = 0.8113$

$B(1) = B(0) = 0$ ,  $B\left(\frac{1}{2}\right) = 1$

2<sup>n</sup> rows in truth table for  $n$  attributes

function DecisionTreeLearning(examples, attributes, parent\_examples) returns tree

if examples is empty then return Plurality-Value(parent.examples)

else if all examples have the same classification then return the classification

else if attributes is empty then return Plurality-Value(examples)

else A  $\leftarrow \arg \max_{a \in \text{attributes}} \text{Importance}(a, \text{examples})$

tree  $\leftarrow$  a new decision tree with root test A

for each value  $v_b$  of A do

  exs  $\leftarrow$  exs  $\text{and } a.v = v_b$

  subtree  $\leftarrow$  DecisionTreeLearning(exs, attributes \ A, examples)

  add a branch to tree with label ( $A = v_b$ ) and subtree subtree

return tree

A: input attribute

V: possible value of input attribute A

Plurality-Value: selects the most common output value among examples.

Reinforcement learning Agent Entity perceiving its env and acting upon it. State Agent observes its state from the environment.

Action Agent acts on the env. through possible actions. Reward indicates how beneficial the last action was w.r.t. solving the task

Policy Function which selects the actions for the agent depending on the current state. Exploration Randomly choosing the next action to create new experience for improving the policy. Exploitation Using the existing policy to decide for the next action.

ROBOTICS

Robots Physical agents that perform tasks by manipulating the physical world. physical quantities

3 Laws of Robotics 1) A robot may not injure a human being or through inaction allow a human being to come to harm. 2) A robot must obey the orders given to it by human beings, except where such orders would conflict with the first law. 3) A robot must protect its own existence as long as such protection does not conflict with the first or second law.

Primary categories Manipulators Robot arms, physically anchored to their workspace. Mobile robots move about their env. using wheels, legs or similar mechanisms. Mobile manipulators Combination of previous categories, e.g. humanoid robots. Sensor Categories Passive sensors Detection of effects generated by other sources in the env (e.g. video/infrared cameras, GPS, light sensor) Active sensors Energy is used to send signals into the env. that are reflected and detected (typically more info provided, but increased power consumption & danger of inference (Sonar, Lidar, Radar)) Proprioceptive sensors inform robot of its own state (angle, force/stain, acceleration sensor) Degree of freedom Number of independent parameters that define the configuration of the robot (body in  $\mathbb{R}^3$  has 6 DOFs: 3 translational and 3 rotational, in  $\mathbb{R}^2$  3 DOFs, 2 translational x/y and 1 rotational  $\theta$ ) Nonholonomic robot has more effective DOFs than controllable DOFs (e.g. car) Holonomic effective DOFs = controllable DOFs Special wheels Omni wheels and Mecanum wheels are used to build holonomic mobile platforms Legged mobile platforms Good in rough terrain. Bad on flat surfaces compared to wheeled platforms Snake robots Versatile (climb, crawl, swim, robust, hard to detect, huge size variations). Very slow. Sensing The detection of a physical presence and its conversion into a signal that can be read by an observer or an instrument. Perception The organization, identification and interpretation of sensory information in order to represent and understand the environment. Difficult due to noisy sensors and partially observable unpredictable and often dynamic env. Localization Problem of finding out where things are, including the robot itself. (Particular instance of filtering / state estimation) Deterministic/stochastic transition model Model state transitions with(out) considering uncertainty  $\rightarrow$  uncertainty because robots are to some degree unpredictable Landmark Stable, recognizable feature of the environment used for navigation. SLAM Simultaneous localization and mapping if no map of the env. is available